



Experiment No. 3
Apply Decision Tree Algorithm on Adult Census Income Dataset and analyze the performance of the model
Date of Performance:
Date of Submission:

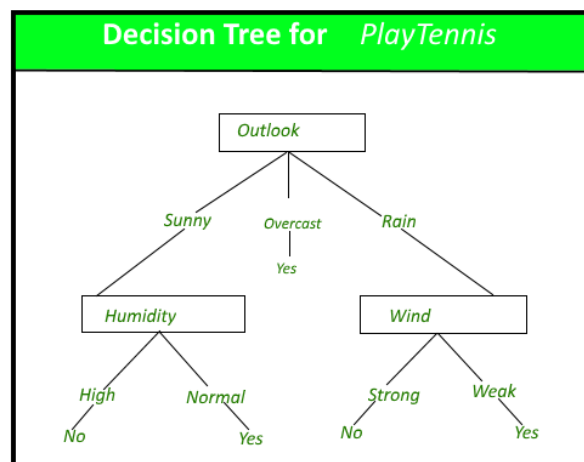


Aim: Apply Decision Tree Algorithm on Adult Census Income Dataset and analyze the performance of the model.

Objective: Able to perform various feature engineering tasks, apply Decision Tree Algorithm on the given dataset and maximize the accuracy, Precision, Recall, F1 score.

Theory:

Decision Tree is the most powerful and popular tool for classification and prediction. A Decision tree is a flowchart-like tree structure, where each internal node denotes a test on an attribute, each branch represents an outcome of the test, and each leaf node (terminal node) holds a class label.



Dataset:

Predict whether income exceeds \$50K/yr based on census data. Also known as "Adult" dataset.

Attribute Information:



Vidyavardhini's College of Engineering & Technology

Department of Computer Engineering

Listing of attributes:

>50K, <=50K.

age: continuous.

workclass: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked.

fnlwgt: continuous.

education: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.

education-num: continuous.

marital-status: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse.

occupation: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces.

relationship: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.

race: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black.

sex: Female, Male.

capital-gain: continuous.

capital-loss: continuous.



hours-per-week: continuous.

native-country: United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad&Tobago, Peru, Hong, Holand-Netherlands.

Code & Output:

```
import os

import numpy as np

import pandas as pd

import matplotlib.pyplot as plt

import seaborn as sns

%matplotlib inline

# Adult dataset path

adult_dataset_path = "/content/adult_dataset.csv"

# Function for loading adult dataset

def load_adult_data(adult_path=adult_dataset_path):

    csv_path = os.path.join(adult_path)

    return pd.read_csv(csv_path)

# Calling load adult function and assigning to a new variable df

df = load_adult_data()

# load top 3 rows values from adult dataset

df.head(3)
```



Vidyavardhini's College of Engineering & Technology

Department of Computer Engineering

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relationship	race	sex	capital.gain	capital.loss	hours.per.week	native.country	income
0	90	?	77053	HS-grad	9	Widowed	?	Not-in-family	White	Female	0	4356	40	United-States	<=50K
1	82	Private	132870	HS-grad	9	Widowed	Exec-managerial	Not-in-family	White	Female	0	4356	18	United-States	<=50K
2	66	?	186061	Some-college	10	Widowed	?	Unmarried	Black	Female	0	4356	40	United-States	<=50K

```
print ("Rows      : " ,df.shape[0])
print ("Columns   : " ,df.shape[1])
print ("\nFeatures : \n" ,df.columns.tolist())
print ("\nMissing values : ", df.isnull().sum().values.sum())
print ("\nUnique values : \n",df.nunique())
```

Rows : 32561

Columns : 15

Features :

```
['age',    'workclass',    'fnlwgt',    'education',    'education.num',
'marital.status',    'occupation',    'relationship',    'race',    'sex',
'capital.gain',    'capital.loss',    'hours.per.week',    'native.country',
'income']
```

Missing values : 0

Unique values :

```
age          73
workclass     9
fnlwgt       21648
education     16
education.num 16
marital.status 7
occupation    15
```

CSL701: Machine Learning Lab



Vidyavardhini's College of Engineering & Technology

Department of Computer Engineering

```
relationship      6
race              5
sex              2
capital.gain     119
capital.loss      92
hours.per.week   94
native.country   42
income           2
```

```
dtype: int64
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   age                   32561 non-null  int64  
 1   workclass              32561 non-null  object  
 2   fnlwgt                 32561 non-null  int64  
 3   education              32561 non-null  object  
 4   education.num          32561 non-null  int64  
 5   marital.status         32561 non-null  object  
 6   occupation             32561 non-null  object  
 7   relationship           32561 non-null  object  
 8   race                   32561 non-null  object  
 9   sex                    32561 non-null  object  
10   capital.gain           32561 non-null  int64  
11   capital.loss           32561 non-null  int64  
12   hours.per.week         32561 non-null  int64  
13   native.country         32561 non-null  object  
14   income                 32561 non-null  object  
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
```

```
df.describe()
```



Vidyavardhini's College of Engineering & Technology

Department of Computer Engineering

	age	fnlwgt	education.num	capital.gain	capital.loss	hours.per.week
count	32561.000000	3.256100e+04	32561.000000	32561.000000	32561.000000	32561.000000
mean	38.581647	1.897784e+05	10.080679	1077.648844	87.303830	40.437456
std	13.640433	1.055500e+05	2.572720	7385.292085	402.960219	12.347429
min	17.000000	1.228500e+04	1.000000	0.000000	0.000000	1.000000
25%	28.000000	1.178270e+05	9.000000	0.000000	0.000000	40.000000
50%	37.000000	1.783560e+05	10.000000	0.000000	0.000000	40.000000
75%	48.000000	2.370510e+05	12.000000	0.000000	0.000000	45.000000
max	90.000000	1.484705e+06	16.000000	99999.000000	4356.000000	99.000000

```
df_check_missing_workclass = (df['workclass']=='?').sum()
```

```
df_check_missing_workclass
```

1836

```
f_check_missing_occupation = (df['occupation']=='?').sum()
```

```
df_check_missing_occupation
```

1843

```
df_missing = (df=='?').sum()
```

```
df_missing
```



	0
age	0
workclass	1836
fnlwgt	0
education	0
education.num	0
marital.status	0
occupation	1843
relationship	0
race	0
sex	0
capital.gain	0
capital.loss	0
hours.per.week	0
native.country	583
income	0
dtype: int64	

```
percent_missing = (df=='?').sum() * 100/len(df)
```

```
percent_missing
```




	0
age	0.000000
workclass	5.638647
fnlwgt	0.000000
education	0.000000
education.num	0.000000
marital.status	0.000000
occupation	5.660146
relationship	0.000000
race	0.000000
sex	0.000000
capital.gain	0.000000
capital.loss	0.000000
hours.per.week	0.000000
native.country	1.790486
income	0.000000
dtype: float64	

df.apply(lambda x: x != '?',axis=1).sum()

	0
age	32561
workclass	30725
fnlwgt	32561
education	32561
education.num	32561
marital.status	32561
occupation	30718
relationship	32561
race	32561
sex	32561
capital.gain	32561
capital.loss	32561
hours.per.week	32561
native.country	31978
income	32561
dtype: int64	



```
# select all categorical variables

df_categorical = df.select_dtypes(include=['object'])

# checking whether any other column contains '?' value

df_categorical.apply(lambda x: x=='?',axis=1).sum()
```

	0
workclass	1836
education	0
marital.status	0
occupation	1843
relationship	0
race	0
sex	0
native.country	583
income	0

dtype: int64

```
# dropping the "?"s from occupation and native.country

df = df[df['occupation'] != '?']

df = df[df['native.country'] != '?']

df.info()
```



```
<class 'pandas.core.frame.DataFrame'>
Index: 30162 entries, 1 to 32560
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   age                   30162 non-null  int64
1   workclass             30162 non-null  object
2   fnlwgt               30162 non-null  int64
3   education             30162 non-null  object
4   education.num         30162 non-null  int64
5   marital.status       30162 non-null  object
6   occupation            30162 non-null  object
7   relationship         30162 non-null  object
8   race                 30162 non-null  object
9   sex                  30162 non-null  object
10  capital.gain          30162 non-null  int64
11  capital.loss          30162 non-null  int64
12  hours.per.week       30162 non-null  int64
13  native.country       30162 non-null  object
14  income               30162 non-null  object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
```

```
from sklearn import preprocessing

# encode categorical variables using label Encoder

# select all categorical variables

df_categorical = df.select_dtypes(include=['object'])

df_categorical.head()
```

	workclass	education	marital.status	occupation	relationship	race	sex	native.country	income
1	Private	HS-grad	Widowed	Exec-managerial	Not-in-family	White	Female	United-States	<=50K
3	Private	7th-8th	Divorced	Machine-op-inspct	Unmarried	White	Female	United-States	<=50K
4	Private	Some-college	Separated	Prof-specialty	Own-child	White	Female	United-States	<=50K
5	Private	HS-grad	Divorced	Other-service	Unmarried	White	Female	United-States	<=50K
6	Private	10th	Separated	Adm-clerical	Unmarried	White	Male	United-States	<=50K

```
# apply label encoder to df_categorical

le = preprocessing.LabelEncoder()

df_categorical = df_categorical.apply(le.fit_transform)

df_categorical.head()

CSL701: Machine Learning Lab
```



Vidyavardhini's College of Engineering & Technology

Department of Computer Engineering

	workclass	education	marital.status	occupation	relationship	race	sex	native.country	income
1	2	11	6	3	1	4	0	38	0
3	2	5	0	6	4	4	0	38	0
4	2	15	5	9	3	4	0	38	0
5	2	11	0	7	4	4	0	38	0
6	2	0	5	0	4	4	1	38	0

```
# first, Drop earlier duplicate columns which had categorical values
```

```
df = df.drop(df_categorical.columns,axis=1)
```

```
df = pd.concat([df,df_categorical],axis=1)
```

```
df.head()
```

	age	fnlwgt	education.num	capital.gain	capital.loss	hours.per.week	workclass	education	marital.status	occupation	relationship	race	sex	native.country	income
1	82	132870	9	0	4356	18	2	11	6	3	1	4	0	38	0
3	54	140359	4	0	3900	40	2	5	0	6	4	4	0	38	0
4	41	264663	10	0	3900	40	2	15	5	9	3	4	0	38	0
5	34	216864	9	0	3770	45	2	11	0	7	4	4	0	38	0
6	38	150601	6	0	3770	40	2	0	5	0	4	4	1	38	0

```
# convert target variable income to categorical
```

```
df['income'] = df['income'].astype('category')
```

```
# Importing train_test_split
```

```
from sklearn.model_selection import train_test_split
```

```
# Putting independent variables/features to X
```

```
X = df.drop('income',axis=1)
```

```
# Putting response/dependent variable/feature to y
```

```
y = df['income']
```

```
X.head(3)
```

	age	fnlwgt	education.num	capital.gain	capital.loss	hours.per.week	workclass	education	marital.status	occupation	relationship	race	sex	native.country	income
1	82	132870	9	0	4356	18	2	11	6	3	1	4	0	38	0
3	54	140359	4	0	3900	40	2	5	0	6	4	4	0	38	0
4	41	264663	10	0	3900	40	2	15	5	9	3	4	0	38	0

```
# Splitting the data into train and test
```



Vidyavardhini's College of Engineering & Technology

Department of Computer Engineering

```
X_train,X_test,y_train,y_test =  
train_test_split(X,y,test_size=0.30,random_state=99)  
  
X_train.head()
```

	age	fnlwgt	education.num	capital.gain	capital.loss	hours.per.week	workclass	education	marital.status	occupation	relationship	race	sex	native.country
24351	42	289636	9	0	0	46	2	11	2	13	0	4	1	38
15626	37	52465	9	0	0	40	1	11	4	7	1	4	1	38
4347	38	125933	14	0	0	40	0	12	2	9	0	4	1	19
23972	44	183829	13	0	0	38	5	9	4	0	1	4	0	38
26843	35	198841	11	0	0	35	2	8	0	12	3	4	1	38

```
# Importing decision tree classifier from sklearn library  
  
from sklearn.tree import DecisionTreeClassifier  
  
# Fitting the decision tree with default hyperparameters, apart from  
# max_depth which is 5 so that we can plot and read the tree.  
  
dt_default = DecisionTreeClassifier(max_depth=5)  
  
dt_default.fit(X_train,y_train)  
  
# Importing classification report and confusion matrix from sklearn  
metrics  
  
from sklearn.metrics import classification_report,confusion_matrix,accuracy_score  
  
# making predictions  
  
y_pred_default = dt_default.predict(X_test)  
  
# Printing classifier report after prediction  
  
print(classification_report(y_test,y_pred_default))
```

	precision	recall	f1-score	support
0	0.86	0.95	0.91	6867
1	0.78	0.52	0.63	2182
accuracy			0.85	9049
macro avg	0.82	0.74	0.77	9049
weighted avg	0.84	0.85	0.84	9049

```
# Printing confusion matrix and accuracy
```



```
print(confusion_matrix(y_test,y_pred_default))
```

```
print(accuracy_score(y_test,y_pred_default))
```

```
[[6553  314]  
 [1039 1143]]  
0.8504807161012267
```

Conclusion:

The model performs well overall with an accuracy of 85%. However, the performance on class 1 (minority class) is not as strong, with lower recall and F1-score. This indicates that the model is better at identifying the majority class (class 0) but struggles with the minority class (class 1).

confusion matrix:

- True Positives (TP) for class 1 = Recall \times Support = $0.52 \times 2182 \approx 1135$
- False Negatives (FN) for class 1 = Support - TP = $2182 - 1135 \approx 1047$
- True Negatives (TN) for class 0 = Recall \times Support = $0.95 \times 6867 \approx 6524$
- False Positives (FP) for class 0 = Support - TN = $6867 - 6524 \approx 343$